

# tradesys: A framework for encoding and backtesting trading systems in R

Robert Sams

July 22, 2009

## 1 Introduction

Design goals: maximum expressibility and tight integration with core R.

Mention other related packages like blotter, TTR, xts, etc...

### 1.1 A formal definition of “trading system”

A *trading system* is an algorithm on a timeseries  $X_t$  that specifies, for each time  $t$ , whether the system’s state is long, short or flat. Mathematically, it is a function  $f(X_t)$  that calculates each state  $s_i \in \{1, 0, -1\}$  on the basis of  $X_1, \dots, X_i$ .  $X_t$  may be as simple as a daily series of closing prices but is often a multivariate series with various price and other data. The states vector combined with the timeseries is the raw material for backtesting research from the calculation of period returns onwards. Let’s call such a combination  $\{X_t, s_t\}$  a *trading system time series*. In this package a trading system time series is represented as class `tsts`.

But what about  $f(X_t)$ , what form does it take?...

### 1.2 Example: Dual-moving average system

```
> library(tradesys)
> library(TTR)
```

This system is long whenever the 60-day moving average of price is above the 120-day moving average and short otherwise. We test it on the S&P 500 index dataset `spx`, which is a zoo matrix and contains daily OHLC and open interest data for about 60 years.

```
> data(spx)
> tail(spx)
```

	Open	High	Low	Close	Volume
2009-05-12	910.52	915.57	896.46	908.35	6871750400
2009-05-13	905.40	905.40	882.80	883.92	7091820000

```

2009-05-14 884.24 898.36 882.52 893.07 6134870000
2009-05-15 892.76 896.97 878.94 882.88 5439720000
2009-05-18 886.07 910.00 886.07 909.71 5702150000
2009-05-19 909.67 916.39 905.22 908.13 6616270000

```

The system can be defined in three lines.

```

> x <- tstts(spx)
> addsig(x, "el", SMA(Close, 60) >= SMA(Close, 120))
> addsig(x, "es", SMA(Close, 60) < SMA(Close, 120))

```

`addsig`—“add signal”—defines the long and short entry criteria, respectively. They take expression objects that must evaluate to logical vectors equal in length to `nrow(x)`. The expressions are evaluated in the normal way using R’s *lazy evaluation* scheme, although `addsig` first puts the columns of `x` into the evaluation frame as named vectors, so `Close` in the above expression evaluates as if it were `data[, ‘Close’]`.

The states vector is now defined and we can do some backtesting.

```

> tail(x)

      Open   High    Low  Close   Volume St
2009-05-12 910.52 915.57 896.46 908.35 6871750400 -1
2009-05-13 905.40 905.40 882.80 883.92 7091820000 -1
2009-05-14 884.24 898.36 882.52 893.07 6134870000 -1
2009-05-15 892.76 896.97 878.94 882.88 5439720000 -1
2009-05-18 886.07 910.00 886.07 909.71 5702150000 -1
2009-05-19 909.67 916.39 905.22 908.13 6616270000 -1

```

The analysis function `equity` is used to calculate period returns and the equity curve.

```

> y <- equity(x, uselog = TRUE)
> tail(y)

```

	Trade	St	Delta	Price	RoR	Equity
2009-05-12	112	-1	1.331220	6.814016	0.018107970	38.02601
2009-05-13	112	-1	1.307543	6.808377	0.007373275	38.30638
2009-05-14	112	-1	1.297973	6.784729	0.030694874	39.48219
2009-05-15	112	-1	1.259318	6.794318	-0.012075942	39.00541
2009-05-18	112	-1	1.274712	6.786796	0.009588169	39.37940
2009-05-19	112	-1	1.262606	6.813082	-0.033188777	38.07244

```

> EquityStats(y[, c("Equity")])

```


RORC	CAGR	ROR%	R2	VOLA	MAXDD
37.07244	0.06317	0.06213	0.84731	0.14157	-0.59800

```

> plot(y[, c("RoR", "Equity")], main = "60/120 DMA -- SPX")

```

### 60/120 DMA -- SPX



Use trades to

enumerate system trades and their holding period returns.

```
> z <- trades(x, uselog = TRUE)
> tail(z)
```

	Phase	ETime	XTime	Time	Numb	EPrice	XPrice	PnL	RoR
107	EL	2006-09-21	2007-09-12	356	244	1324.89	1471.10	146.21	0.264117052
108	ES	2007-09-12	2007-11-09	58	42	1471.10	1467.59	3.51	0.006027153
109	EL	2007-11-09	2008-01-03	55	36	1467.59	1447.55	-20.04	-0.034689959
110	ES	2008-01-03	2008-06-10	159	109	1447.55	1358.98	88.57	0.159301490
111	EL	2008-06-10	2008-07-21	41	28	1358.98	1261.82	-97.16	-0.187159273
112	ES	2008-07-21	2009-05-19	302	209	1261.82	909.67	352.15	0.825619186

What prices are assumed in these calculations, as `spx` contains columns for open, high, low and close prices? By default, the prices used in these performance calculations is the left-most column, which is this case, is the open price.

```
> unlist(pricecols(x))
```

valuation	enterlong	entershort	exitlong	exitshort	rolllong	rollshort
1	1	1	1	1	1	1

This won't due. We can't calculate our signal on Monday's closing price whilst trading on that signal using Monday's open price!

The `'tsts'` class has a very flexible mechanism for defining the price context of a trading system timeseries.

```
> pricecols(x) <- "Close"
> unlist(pricecols(x))
```

```
valuation  enterlong entershort  exitlong  exitshort  rolldlong  rolldshort
"Close"    "Close"    "Close"    "Close"    "Close"    "Close"    "Close"
```

This specifies that the system will assume that all trades are executed at the closing price. This is an improvement. But let's say that we want the system to compute signals on closing prices, position valuations at closing prices, and trades on the *following day's* open price. This is a good place to introduce to the `addcol` function, for it is our solution to this problem.

```
> addcol(x, "Next", c(embed(Open, 2)[, 1], NA))
> tail(x)
```

	Open	High	Low	Close	Volume	Next	St
2009-05-12	910.52	915.57	896.46	908.35	6871750400	905.40	-1
2009-05-13	905.40	905.40	882.80	883.92	7091820000	884.24	-1
2009-05-14	884.24	898.36	882.52	893.07	6134870000	892.76	-1
2009-05-15	892.76	896.97	878.94	882.88	5439720000	886.07	-1
2009-05-18	886.07	910.00	886.07	909.71	5702150000	909.67	-1
2009-05-19	909.67	916.39	905.22	908.13	6616270000	NA	-1

So now `x` has an additional column named "Next", which contains the open price of the *next* day. Like the signal expressions, the expression that defines this column is stored in `x` so that the column is recalculated everytime new data is assigned to `x` via `coredata`. Now we can update our pricing assumptions with another call to `pricecols`.

```
> pricecols(x) <- list(valuation = "Close", entry = "Next", exit = "Next")
```

Let's Re-run `trades` and examine an arbitrary example, say trade 109.

```
> trades(x)[109, ]
```

	Phase	ETime	XTime	Time	Numb	EPrice	XPrice	PnL	RoR
109	EL	2007-11-09	2008-01-03	55	36	1467.59	1444.01	-23.58	NA

You can confirm for yourself that the trade entry on 11-09 is done at the Open price on 11-12...

```
> window(spx, start = "2007-11-09", end = "2007-11-12")
```

	Open	High	Low	Close	Volume
2007-11-09	1467.59	1474.09	1448.51	1453.70	4587050000
2007-11-12	1453.66	1464.94	1438.53	1439.18	4192520000

```
> window(spx, start = "2008-01-03", end = "2008-01-04")
```

	Open	High	Low	Close	Volume
2008-01-03	1447.55	1456.80	1443.73	1447.16	3429500000
2008-01-04	1444.01	1444.01	1411.19	1411.63	4166000000

## 2 Splicing Timeseries

### Computational details

The results in this paper were obtained using R 2.8.0 with the packages `tradesys` 0.1 and `zoo` 1.5-4 R itself and all packages used are available from CRAN at <http://CRAN.R-project.org/>.